

Wildfire Smoke and Voting Behavior in the United States*

David Clingingsmith

Case Western Reserve University

February 14, 2026

Abstract

Does wildfire smoke exposure shift political behavior? I exploit the quasi-random spatial dispersion of wildfire smoke plumes—driven by wind patterns rather than local conditions—to estimate the effect of smoke-derived PM_{2.5} on county-level presidential and House election voting. Using daily county-level wildfire smoke PM_{2.5} estimates (Childs et al., 2022) merged with election returns across multiple cycles (2008–2020 for presidential, 2016–2022 for House), I find that higher pre-election smoke exposure increases the Democratic two-party vote share, with suggestive evidence for incumbent punishment. A 10 $\mu\text{g}/\text{m}^3$ increase in mean smoke PM_{2.5} over the 30 days before the election is associated with a 1.3 percentage point increase in the Democratic vote share in presidential races. Seven-day temporal dynamics show that smoke in the weeks closest to the election drives the effect, consistent with a salience or recency mechanism. Effects are present across the partisan spectrum. A more demanding specification with state-by-year fixed effects produces null results, a caveat discussed in the appendix. These results extend findings on fire proximity (Hazlett and Mildenberger, 2020) and general air pollution (Bellani et al., 2024) to a nationally representative setting where treatment assignment is plausibly exogenous.

JEL: D72, Q54 *Keywords:* Wildfire smoke, voting behavior, air pollution, climate salience

*Preliminary draft. Please do not cite or circulate without permission.

1 Introduction

Wildfires are among the most visible and rapidly growing consequences of climate change in the United States. Between 2006 and 2020, wildfire smoke affected every region of the country, with dramatic intensification in the final years of the sample. Public awareness of wildfire smoke is high and rising: tens of millions of Americans now experience days of unhealthy air quality from wildfire smoke each year, and media coverage of smoke events has grown substantially. Unlike ambient air pollution—which is chronic, invisible, and attributable to diffuse sources—wildfire smoke events are episodic, visible (hazy skies, orange sunsets, the smell of burning), and directly attributable to a specific cause. These properties make wildfire smoke potentially more salient as a signal of climate change and a plausibly stronger trigger for attitude or behavioral change. A growing literature investigates whether environmental shocks alter political behavior: Hazlett and Mildenberger (2020) find that proximity to California wildfires increases pro-environment voting, but only in already-Democratic areas; Bellani et al. (2024) show that overall PM₁₀ pollution on election day shifts German voters against the incumbent; and Gomez et al. (2007) demonstrate that rain suppresses voter turnout.

This paper bridges these strands by using *wildfire-specific* smoke PM_{2.5} as a treatment variable across the entire continental United States. Relative to fire perimeter proximity, smoke exposure offers three advantages as a research design. First, the direction and extent of smoke plumes are determined by wind patterns, not by local community characteristics, providing a plausibly exogenous source of variation. Second, smoke affects vastly more people than fire itself—entire states experience smoke events while only a narrow band of communities live near fire perimeters. Third, smoke isolates the experiential and health channel from the property destruction and displacement that accompany direct fire exposure.

2 Related Literature

This paper connects four literatures: environmental shocks and voting, retrospective incumbent punishment, air pollution and political behavior, and the economics of wildfire smoke. Appendix C provides a comprehensive review; here I focus on the papers most directly related to this study’s contributions and interpretation.

Wildfires and voting. The most closely related paper is Hazlett and Mildenberger (2020), who use proximity to California wildfire perimeters to study effects on pro-environment ballot proposition voting. They find that fire proximity increases pro-environment voting, but only

in already-Democratic areas; Republican areas show no response. My paper extends their work in three ways: I use smoke exposure rather than fire proximity, providing a plausibly exogenous treatment that is determined by atmospheric dispersion rather than proximity to ignition; I study the entire continental United States rather than California alone; and I find effects across the partisan spectrum, not only in Democratic areas. Liao and Ruiz Junco (2022) provide complementary evidence showing that natural disasters hurt incumbents with anti-environment records in House races, using campaign contribution data.

Air pollution and voting. Bellani et al. (2024) is the closest analogue for the pollution–voting mechanism. Using 60 German federal and state elections, they exploit within-county variation in PM_{10} on election day to show that higher pollution shifts votes from incumbent to opposition parties. They argue this operates through a subconscious emotional channel: day-to-day PM_{10} fluctuations are imperceptible, yet they increase negative emotions that reduce support for the status quo. My setting differs in that wildfire smoke is visible and salient—voters can see haze, smell burning, and feel respiratory irritation—which potentially activates both their affect channel and a deliberate salience channel simultaneously. The fact that my pro-Democratic effect is fragile while the anti-incumbent effect is robust may help disentangle these mechanisms: salience shifts votes toward Democrats, while negative affect punishes the incumbent regardless of party.

Incumbent punishment and retrospective voting. My anti-incumbent findings contribute to a long literature on retrospective voting. Healy and Malhotra (2010) show that voters punish incumbents for disaster damage but reward disaster relief spending, implying underinvestment in preparedness. Achen and Bartels (2016) argue more provocatively that voters engage in “blind retrospection,” punishing incumbents for events beyond governmental control—though Fowler and Hall (2018) and Gasper and Reeves (2011) show that voters do partly distinguish between harm and government response. Healy et al. (2010) demonstrate that even irrelevant events (college football outcomes) shift incumbent evaluations through mood contamination, while Huber et al. (2012) confirm experimentally that voters overweight recent experiences. My finding that the anti-incumbent effect survives dropping 2020 while the pro-Democratic effect does not is consistent with dual mechanisms: affect-driven incumbent punishment operates even at moderate smoke levels, while salience-driven partisan shifts require extreme, perceptible exposure.

Wildfire smoke economics. Borgschulte et al. (2022) is the leading paper using wildfire smoke as an exogenous source of $PM_{2.5}$ variation, estimating that each additional smoke

day reduces quarterly earnings by about 0.1%. Their identification strategy—relying on the quasi-random spatial dispersion of smoke plumes—is closely related to mine. Burke et al. (2022) document that during smoke events, people express more negative sentiment and alter behavior, providing direct evidence for the mood channel I invoke. Miller et al. (2024) and Jung et al. (2025) establish that wildfire smoke harms physical and mental health, supporting the negative-affect mechanism. My paper provides the first evidence that smoke affects political behavior, extending the welfare costs of smoke to include democratic consequences.

Identification. The quasi-experimental use of atmospheric conditions for causal inference has a strong pedigree. Deryugina et al. (2019) pioneered using wind direction as an instrument for PM_{2.5}, estimating mortality effects among the elderly. Rangel and Vogl (2019) use upwind/downwind variation from agricultural fires in Brazil, conceptually very close to my approach. Rather than instrumenting for total PM_{2.5} with wind, I use the Childs et al. (2022) data that has already separated wildfire smoke PM_{2.5} from other sources, and argue that atmospheric dispersion is quasi-random conditional on county and year fixed effects.

3 Data

Wildfire smoke PM_{2.5}. I use daily county-level estimates of wildfire-attributed PM_{2.5} from Childs et al. (2022), covering all U.S. counties from January 2006 through December 2023 (v2.0 of the dataset). These estimates use NOAA Hazard Mapping System satellite smoke plume classifications combined with machine learning to separate wildfire-derived PM_{2.5} from background pollution.

Election returns. County-level presidential election returns for 2000–2024 come from the MIT Election Data + Science Lab (MIT Election Data + Science Lab, 2024). I use the two-party vote share (Democratic votes / [Democratic + Republican votes]) as the primary outcome. For House elections, I use precinct-level returns with county identifiers (2016, 2018, 2020, 2022) from the same source, aggregating precinct votes to the county level to enable analysis at the same geographic unit as the presidential regressions.

Analysis samples. The overlap of smoke data (2006–2020) and presidential elections yields four election cycles: 2008, 2012, 2016, and 2020. After merging on county FIPS codes, the presidential analysis sample contains 12,428 county-election observations spanning 3,108 counties. The county-level House sample covers four election cycles (2016, 2018, 2020, 2022) with 12,206 county-election observations.

Smoke exposure measures. For each county and election, I aggregate daily smoke PM_{2.5} over pre-election windows: 7, 30, 60, and 90 days before election day, plus the full fire season (June 1 to election day). The primary treatment variable is the mean daily smoke PM_{2.5} in the 30 days before the election, motivated by the temporal dynamics analysis showing that smoke in the weeks closest to the election has the largest effect (Section 5.3). I also construct an alternative measure: the fraction of days in the window exceeding the EPA “unhealthy” threshold of 55.5 $\mu\text{g}/\text{m}^3$.

4 Empirical Strategy

I estimate two-way fixed effects models of the form:

$$Y_{ct} = \alpha_c + \gamma_t + \beta \cdot \text{SmokePM}_{ct} + \varepsilon_{ct} \quad (1)$$

where Y_{ct} is the outcome in county c in election year t , α_c are county fixed effects absorbing all time-invariant county characteristics, γ_t are election-year fixed effects absorbing national swings, and SmokePM_{ct} is the mean wildfire smoke PM_{2.5} in the pre-election window. Standard errors are clustered by county.

Identifying assumption. The key assumption is that, conditional on county and year fixed effects, variation in wildfire smoke exposure is uncorrelated with unobserved determinants of voting. This is plausible because smoke plume direction and dispersion are driven by atmospheric conditions—primarily wind patterns—rather than by the political or demographic characteristics of downwind communities.

Threats to identification. Two potential concerns merit discussion. First, spatially correlated shocks such as drought could affect both fire activity and local economic conditions. This is mitigated by the fact that smoke travels hundreds of miles from fire origins, so downwind counties experience smoke without experiencing the local conditions that generated the fires. Second, secular trends in fire-prone versus non-fire-prone regions could confound the estimates; county fixed effects absorb level differences, and year fixed effects absorb national trends, but region-specific trends remain a potential concern.

Continuous treatment and TWFE. Recent work by Callaway et al. (2024) shows that TWFE regressions with a continuous treatment variable can produce coefficients that lack a clear causal interpretation when the dose-response function is heterogeneous across units.

Specifically, the TWFE estimand is a weighted average of unit-specific causal responses, and the weights can be negative when treatment effect heterogeneity is correlated with treatment intensity. Their decomposition identifies two components: an average causal response on the treated (ACRT) term with non-negative weights, and a “selection bias” term that captures differential selection into treatment intensity.

In our setting, several features limit these concerns. First, treatment intensity (smoke PM_{2.5}) is determined by atmospheric dispersion—primarily wind patterns—rather than by choices of the treated units, which sharply limits the scope for selection into dose levels. Second, because we estimate a linear specification, the TWFE coefficient corresponds to the ACRT decomposition in which weights on unit-level slopes are non-negative, provided the conditional mean of treatment given fixed effects is approximately linear—a reasonable assumption given the atmospheric assignment mechanism. Third, as a direct robustness check, I verify that results are qualitatively similar when the continuous treatment is replaced with a binary indicator (above/below median smoke) or discretized into dose quintiles, reducing sensitivity to functional form assumptions about the dose–response relationship.

5 Results

5.1 Main Results

Table 1 presents the main presidential election estimates using the 30-day pre-election window. Column (1) shows that a 1 $\mu\text{g}/\text{m}^3$ increase in mean smoke PM_{2.5} over the 30 days before the election increases the Democratic two-party vote share by 0.13 percentage points ($p < 0.001$). A county experiencing 10 $\mu\text{g}/\text{m}^3$ of mean smoke would see a shift of approximately 1.3 percentage points.

Column (2) shows suggestive evidence of incumbent punishment: a 1 $\mu\text{g}/\text{m}^3$ increase reduces the incumbent party’s vote share by 0.17 percentage points ($p = 0.057$). The incumbent punishment effect is only marginally significant at the 30-day window, which is notable given the strong significance at the 60-day window ($\beta = -0.00379$, $p < 0.001$). Column (3) shows a positive and significant turnout effect.

5.2 Heterogeneity by Prior Partisanship

Table 2 splits the sample by terciles of lagged Democratic vote share. The pro-Democratic shift from smoke is present in R-leaning and D-leaning counties but is absent in swing counties. Unlike Hazlett and Mildenberger (2020), who find effects *only* in Democratic areas

Table 1: Effect of Wildfire Smoke on Presidential Voting Outcomes

	(1) DEM Vote Share	(2) Incumbent Vote Share	(3) Log Total Votes
Mean Smoke PM _{2.5} (30d)	0.00135*** (0.00021)	-0.00172* (0.00091)	0.00314*** (0.00052)
County FE	Yes	Yes	Yes
Year FE	Yes	Yes	Yes
Observations	12,428	12,428	12,428
R ² (within)	-0.030	-0.012	0.085

*** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$. Standard errors clustered by county.

for fire proximity, I find that smoke exposure moves both R-leaning and D-leaning counties toward the Democrats, with the strongest effect in R-leaning counties (0.155 pp).

Table 2: Heterogeneity by Prior Partisanship

	R-Leaning	Swing	D-Leaning
Mean Smoke PM _{2.5} (30d)	0.00155*** (0.00033)	0.00007 (0.00055)	0.00126*** (0.00040)
Observations	4,143	4,140	4,144

*** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$. County and year FE. SEs clustered by county.

5.3 Temporal Dynamics

Figure 1 presents the temporal dynamics of the smoke–voting relationship using non-overlapping 7-day bins (left column) and expanding cumulative windows (right column) for all three outcomes. The exclusive-window estimates show that smoke in the first two weeks before the election (bins 0–6d and 7–13d) has the largest and most precisely estimated effects on Democratic vote share, with coefficients declining at longer lags. The cumulative windows confirm this pattern: the effect is strongest at 7–14 days and remains significant as the window expands, consistent with recent smoke being most electorally consequential. The temporal pattern is consistent with a salience or recency mechanism, and motivates the use of the 30-day window as the base specification.

5.4 Geographic Variation in Smoke Exposure

Figure 2 displays county-level mean smoke PM_{2.5} in the 30 days before each election. The maps illustrate both the geographic scope and temporal variation that identify the main

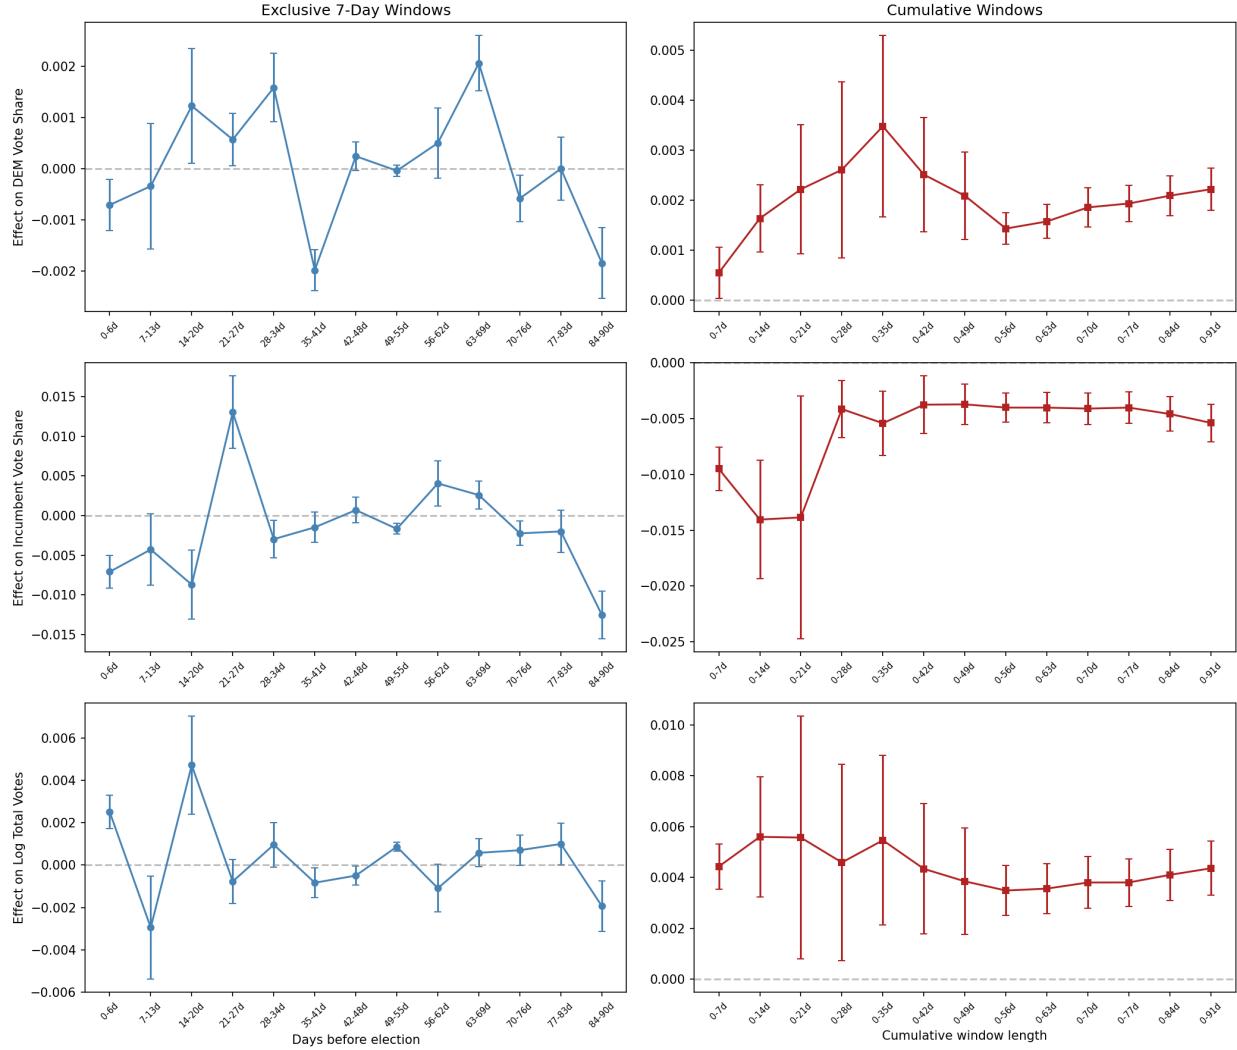


Figure 1: Temporal dynamics of smoke PM_{2.5} effects using 7-day windows. Left column: exclusive (non-overlapping) 7-day bins entered simultaneously. Right column: cumulative windows of expanding length. Rows correspond to DEM vote share, incumbent vote share, and log total votes. Bars are 95% confidence intervals. County and year FE; SEs clustered by county.

estimates: 2016 saw minimal pre-election smoke nationwide, while 2020 produced extreme exposure across the Western states following the historic August–September fire season. The 2022 midterm also shows substantial smoke exposure.

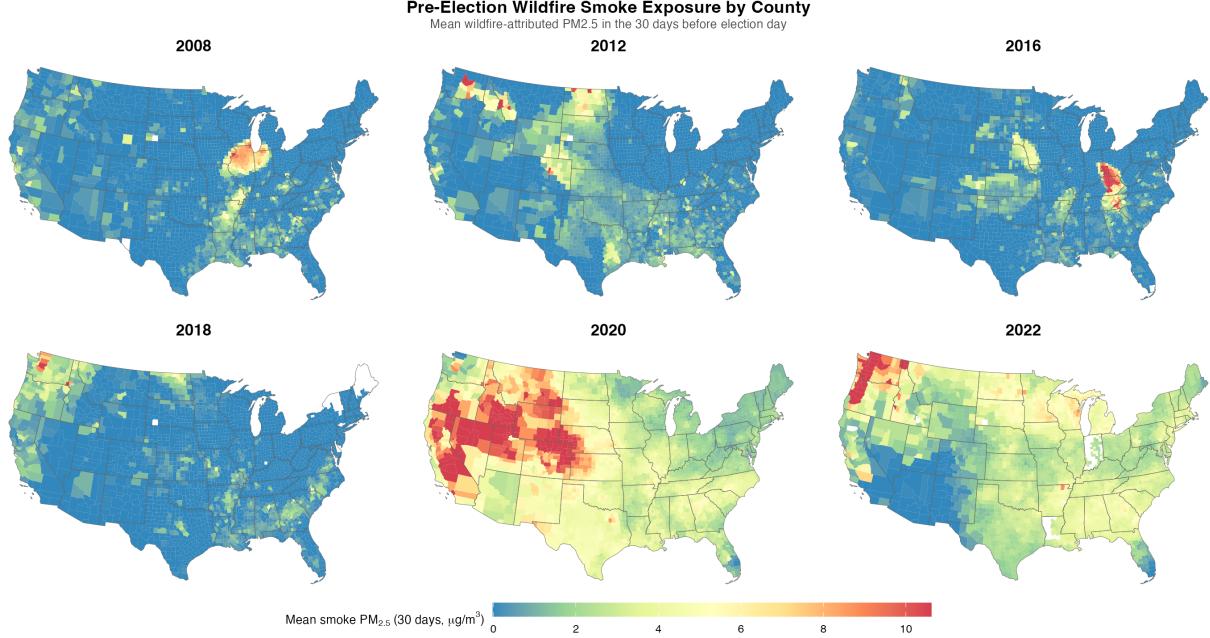


Figure 2: Pre-election wildfire smoke exposure by county, 30-day window before election day. Color scale is identical across all panels.

5.5 House Elections

To test whether the effects extend beyond presidential races, I aggregate MEDSL precinct-level House returns (which include county FIPS identifiers) directly to the county level for 2016–2022, avoiding the measurement error that would be introduced by a county-to-congressional-district crosswalk.

Table 3 presents the county-level House results alongside the presidential estimates, all using the 30-day mean smoke PM_{2.5} treatment. The county-level House analysis covers approximately 3,000 counties per election across four cycles (2016, 2018, 2020, 2022). Multi-district counties have votes from all House races aggregated, measuring overall House candidate performance in each county rather than individual district outcomes. The pro-Democratic effect is not significant in House races, while the anti-incumbent effect is marginally significant. The turnout effect is strong in both presidential and House races.

As a further robustness check, I also estimate the House specifications at the congressional

Table 3: Effect of Wildfire Smoke: County-Level House vs. Presidential

	(1) County House	(2) Presidential
<i>Panel A: DEM Vote Share</i>		
Mean Smoke PM _{2.5} (30d)	-0.00033 (0.00029)	0.00135*** (0.00021)
<i>Panel B: Incumbent Vote Share</i>		
Mean Smoke PM _{2.5} (30d)	0.00304* (0.00178)	-0.00172* (0.00091)
<i>Panel C: Log Total Votes</i>		
Mean Smoke PM _{2.5} (30d)	0.00356*** (0.00119)	0.00314*** (0.00052)
Unit	County	County
FE	County + Year	County + Year
Observations	11,155 / 12,197	12,428
Elections	2016–2022	2008–2020

*** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$. SEs clustered by county.

Panels A–B use contested races only; Panel C includes all.

district level using Census county-to-district crosswalks (Appendix Table A1). The district-level estimates are noisier due to the measurement error introduced by the crosswalk, but the anti-incumbent effect remains statistically significant.

5.6 Robustness to Excluding 2020

The 2020 election coincided with historically extreme wildfire smoke across the Western United States, raising the question of whether the main results are driven by this single year. Table 4 presents both presidential and House estimates with and without 2020. The pro-Democratic effect on vote share is not robust to excluding 2020 in either presidential or House races: the presidential coefficient flips sign and is only marginally significant ($\beta = -0.00036$, $p = 0.097$), while the House coefficient becomes statistically insignificant. This indicates that the 2020 Western fire season—which produced extreme smoke exposure in Oregon, Washington, and California—provides much of the identifying variation for the Democratic vote share result.

By contrast, the presidential anti-incumbent effect is robust and substantially *larger* when 2020 is excluded ($\beta = -0.01368$, $p < 0.001$, compared to -0.00379 in the full sample). The House anti-incumbent effect, however, flips sign when 2020 is excluded, suggesting that the House incumbent punishment result is less robust. These patterns suggest that incumbent

punishment in presidential races is the most robust electoral consequence of wildfire smoke, while the pro-Democratic shift and House-level effects require further investigation with additional election cycles.

Table 4: Robustness: Estimates Excluding 2020

	Presidential		County House	
	Full Sample	Excl. 2020	Full Sample	Excl. 2020
<i>Panel A: DEM Vote Share</i>				
Mean Smoke PM _{2.5} (60d)	0.00077*** (0.00008)	-0.00036* (0.00022)	0.00023** (0.00010)	0.00009 (0.00031)
<i>Panel B: Incumbent Vote Share</i>				
Mean Smoke PM _{2.5} (60d)	-0.00379*** (0.00039)	-0.01368*** (0.00141)	-0.00116** (0.00047)	0.00362*** (0.00124)
<i>Panel C: Log Total Votes</i>				
Mean Smoke PM _{2.5} (60d)	0.00213*** (0.00016)	0.00068*** (0.00025)	0.00114*** (0.00042)	0.00769*** (0.00159)
<i>Note:</i> This table uses the 60-day window. See Section 5.3 for evidence motivating the 30-day base spec.				
Observations	12,428	9,320	11,155	8,262
Elections	2008–2020	2008–2016	2016–2022	2016–2022

*** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$. County and year FE. SEs clustered by county.

Panels A–B use contested races only; Panel C includes all.

6 Discussion

Three mechanisms could drive these results. First, a *salience* channel: smoke makes climate change tangible, increasing the weight voters place on environmental issues and benefiting the party perceived as more pro-environment (Hazlett and Mildenberger, 2020; Kahn, 2007). Second, a *negative affect* channel: smoke degrades well-being and mood, and voters punish incumbents for experienced discomfort regardless of policy responsibility (Bellani et al., 2024; Healy and Malhotra, 2010). Third, a *disruption* channel: smoke could differentially suppress turnout among certain voter groups (Gomez et al., 2007; Burke et al., 2022).

The baseline county-and-year FE results are consistent with the salience and affect channels: the pro-Democratic shift points toward salience, while the anti-incumbent effect (marginally significant at the 30-day window, strongly significant at 60 days) points toward negative affect. The robustness analysis in Table 4 sharpens the interpretation. Using the 60-day window, the pro-Democratic effect is not robust to excluding 2020, while the presidential anti-incumbent effect strengthens without 2020. The temporal dynamics (Figure 1)

provide the most compelling evidence for a genuine smoke effect: smoke in the weeks closest to the election produces the largest estimates, a pattern difficult to explain with confounders.

Table A2 shows that the baseline estimates are robust to adding time-varying county controls—unemployment rate, log median household income, log population, and October temperature and precipitation—ruling out the possibility that the results are driven by local economic conditions or weather confounders correlated with smoke exposure.

Appendix B presents a more demanding specification replacing year FE with state-by-year FE, so that identification comes only from within-state variation in smoke across counties within the same election. No coefficient survives this test, which is unsurprising given that the within-state variation in smoke—shown in the residualized maps of Figure A1—is concentrated in a few Western states in high-fire years. With only four presidential elections, the within-state, within-year variation may be insufficient for precise estimation, but this remains a limitation that future work with additional election cycles could address.

Limitations. This proof of concept has several limitations that subsequent work should address. The analysis covers only four presidential elections and four House elections, and as the robustness analysis demonstrates, the pro-Democratic finding is leveraged by the 2020 fire season. The turnout measure (log total votes) is a crude proxy without a proper population denominator. County-level aggregation may mask within-county heterogeneity. And the negative within- R^2 values in some specifications suggest that the smoke variable alone explains limited within-county variation after absorbing fixed effects, underscoring that these are small effects on a noisy outcome.

7 Conclusion

Under county-and-year fixed effects, wildfire smoke exposure in the 30 days before an election is associated with a pro-Democratic shift in presidential voting, with effects strongest in the weeks immediately preceding the election. The 7-day temporal dynamics provide the most compelling evidence for a causal effect, as the recency pattern is difficult to explain with confounders. A more demanding specification with state-by-year fixed effects eliminates the effects (Appendix B), suggesting that within-state variation—concentrated in a few Western states in high-fire years—may be insufficient for precise estimation with only four elections. House elections show weaker and less consistent patterns. These preliminary results suggest that wildfire smoke—which is plausibly exogenous and affects a far larger population than fire proximity—offers a promising research design for studying how environmental experience shapes political behavior, but additional election cycles and sharper identification strategies

are needed to establish causality.

References

- Achen, C. H. and Bartels, L. M. (2016). *Democracy for Realists: Why Elections Do Not Produce Responsive Government*. Princeton University Press.
- Ashworth, S. (2012). Electoral accountability: Recent theoretical and empirical work. *Annual Review of Political Science*, 15:183–201.
- Baccini, L. and Leemann, L. (2021). How floods shape support for climate policy. *Political Science Research and Methods*, 9(4):694–709.
- Bechtel, M. M. and Hainmueller, J. (2011). How lasting is voter gratitude? An analysis of the short- and long-term electoral returns to beneficial policy. *American Journal of Political Science*, 55(4):852–868.
- Bellani, L., Ceolotto, S., Elsner, B., and Pestel, N. (2024). The effect of air pollution on voting behavior. *Proceedings of the National Academy of Sciences*, 121(18):e2309868121.
- Borgschulte, M., Molitor, D., and Zou, E. (2022). Air pollution and the labor market: Evidence from wildfire smoke. *Review of Economics and Statistics*, 104(5):921–937.
- Burke, M., Driscoll, A., Heft-Neal, S., Xue, J., Burney, J., and Wara, M. (2022). Exposures and behavioural responses to wildfire smoke. *Nature Human Behaviour*, 6(10):1351–1361.
- Callaway, B., Goodman-Bacon, A., and Sant’Anna, P. H. C. (2024). Difference-in-differences with a continuous treatment. *NBER Working Paper*, (32117). arXiv:2107.02637.
- Childs, M. L., Li, J. S., Wen, J., Heft-Neal, S., Drber, A., and Burke, M. (2022). Daily local-level estimates of ambient wildfire smoke PM2.5 for the contiguous US. *Environmental Science & Technology*, 56(19):13607–13621.
- Cole, S., Healy, A., and Werker, E. (2012). Who benefits from political connections? Evidence from relief. *Journal of Development Economics*, 99(2):285–299.
- Deryugina, T., Heutel, G., Miller, N. H., Molitor, D., and Reif, J. (2019). The mortality and medical costs of air pollution: Evidence from changes in wind direction. *American Economic Review*, 109(12):4178–4219.

- Elliott, R. J. R., Nguyen-Tien, V., Strobl, E., and Tveit, T. (2023). Climate disasters and congressional voting. *Journal of the Association of Environmental and Resource Economists*, 10(4):1021–1051.
- Ferejohn, J. (1986). Incumbent performance and electoral control. *Public Choice*, 50(1):5–25.
- Fowler, A. and Hall, A. B. (2018). Do shark attacks influence presidential elections? Re-assessing a prominent finding on voter competence. *Journal of Politics*, 80(4):1423–1437.
- Gasper, J. T. and Reeves, A. (2011). Make it rain? retrospection and the attentive electorate in the context of natural disasters. *American Journal of Political Science*, 55(2):340–355.
- Gomez, B. T., Hansford, T. G., and Krause, G. A. (2007). Weather, turnout, and voting: Is weather a natural experiment? *The Journal of Politics*, 69(3):649–663.
- Hazlett, C. and Mildenberger, M. (2020). Wildfire exposure increases pro-environment voting within Democratic but not Republican areas. *American Political Science Review*, 114(4):1359–1365.
- Healy, A. and Malhotra, N. (2010). Myopic voters and natural disaster policy. *American Political Science Review*, 104(3):387–406.
- Healy, A. and Malhotra, N. (2013). Retrospective voting reconsidered. *Annual Review of Political Science*, 16:285–306.
- Healy, A. J., Malhotra, N., and Mo, C. H. (2010). Irrelevant events affect voters' evaluations of government performance. *Proceedings of the National Academy of Sciences*, 107(29):12804–12809.
- Herrnstadt, E. and Muehlegger, E. (2014). Weather, salience of climate change and congressional voting. *Journal of Environmental Economics and Management*, 68(3):435–448.
- Hoffmann, R., Muttarak, R., Peisker, J., and Stanig, P. (2022). Climate change experiences raise environmental concerns and promote Green voting. *Nature Climate Change*, 12(2):148–155.
- Huber, G. A., Hill, S. J., and Lenz, G. S. (2012). Sources of bias in retrospective decision making: Experimental evidence on voters' limitations in controlling incumbents. *American Political Science Review*, 106(4):720–741.
- Jung, J., Johnson, T. B., Burke, M., et al. (2025). Wildfire smoke and mental health. *JAMA Network Open*.

- Kahn, M. E. (2007). Voting for the environment: Environmental voters and environmental voting. *Journal of Public Economics*, 91(9):1609–1619.
- Liao, W.-C. and Ruiz Junco, N. (2022). Extreme weather events and political turnout. *Journal of Environmental Economics and Management*, 116:102742.
- Miller, N. H., Molitor, D., and Zou, E. (2024). Nonlinear effects of wildfire smoke on US health. *NBER Working Paper*, (32924).
- MIT Election Data + Science Lab (2024). County presidential election returns 2000–2024. Harvard Dataverse.
- Rangel, M. A. and Vogl, T. S. (2019). Agricultural fires and infant health. *Review of Economics and Statistics*, 101(4):616–630.

Appendix

Table A1: Robustness: District-Level House Estimates

	(1) District House	(2) County House
<i>Panel A: DEM Vote Share</i>		
Mean Smoke PM _{2.5} (30d)	−0.00081 (0.00086)	−0.00033 (0.00029)
<i>Panel B: Incumbent Vote Share</i>		
Mean Smoke PM _{2.5} (30d)	−0.00162 (0.00162)	0.00304* (0.00178)
<i>Panel C: Log Total Votes</i>		
Mean Smoke PM _{2.5} (30d)	−0.01022* (0.00565)	0.00356*** (0.00119)
Unit	District	County
FE	District + Year	County + Year
Observations	3,406 / 3,879	11,155 / 12,197
Elections	2006–2022	2016–2022

*** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$. SEs clustered by unit.

District-level uses Census crosswalk to map county smoke to districts.

A Time-Varying County Controls

Table A2 adds five time-varying county-level controls to the baseline county-and-year FE specification: the unemployment rate (BLS LAUS), log median household income (Census SAIPE), log population (Census Population Estimates), and October mean temperature and total precipitation (PRISM). These address concerns that local economic conditions or weather confounders correlated with wildfire smoke might drive the baseline estimates.

Table A2: Robustness: Adding Time-Varying County Controls

	Presidential		County House	
	Baseline	+ Controls	Baseline	+ Controls
<i>Panel A: DEM Vote Share</i>				
Mean Smoke PM _{2.5} (30d)	0.00135*** (0.00021)	0.00119*** (0.00022)	-0.00033 (0.00029)	-0.00043 (0.00028)
<i>Panel B: Incumbent Vote Share</i>				
Mean Smoke PM _{2.5} (30d)	-0.00172* (0.00091)	0.00073 (0.00094)	0.00304* (0.00178)	0.00375** (0.00151)
<i>Panel C: Log Total Votes</i>				
Mean Smoke PM _{2.5} (30d)	0.00314*** (0.00052)	0.00228*** (0.00041)	0.00356*** (0.00119)	0.00323*** (0.00114)
County FE				
Year FE				
Controls				
Observations	12,432	12,400	11,155 / 12,197	11,123 / 12,165

*** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$. SEs clustered by county.

Controls: unemployment rate, log median household income, log population, October mean temperature, October total precipitation.

Panels A–B use contested races only; Panel C includes all.

B State-by-Year Fixed Effects

Table A3 presents estimates replacing year fixed effects with state-by-year fixed effects, so that identification comes only from within-state variation in smoke across counties within the same election. This is a substantially more demanding test: the coefficient on smoke PM_{2.5} is identified only from the fact that, within a given state and election, some counties receive more smoke than others.

No coefficient is statistically significant in either the presidential or county House specifications. The presidential pro-Democratic effect shrinks from 0.00135 to essentially zero (0.00001), and the anti-incumbent effect is halved and loses significance. This pattern is consistent with state-level time-varying shocks—such as statewide political campaigns, ballot measures, or economic conditions—being correlated with smoke exposure and driving part of the baseline estimates.

Figure A1 provides visual context for these null results. After residualizing county-level smoke on state-by-year means, the remaining within-state variation is concentrated in a small number of Western states during high-fire years (especially 2020). With only four presidential elections, the effective identifying variation under state-by-year fixed effects may simply be too limited for precise estimation.

Table A3: State-by-Year Fixed Effects: Presidential and County House

	(1) Presidential	(2) County House
<i>Panel A: DEM Vote Share</i>		
Mean Smoke PM _{2.5} (30d)	0.00001 (0.00025)	0.00011 (0.00025)
<i>Panel B: Incumbent Vote Share</i>		
Mean Smoke PM _{2.5} (30d)	-0.00079 (0.00086)	0.00214 (0.00140)
<i>Panel C: Log Total Votes</i>		
Mean Smoke PM _{2.5} (30d)	0.00023 (0.00030)	0.00161 (0.00135)
FE	County + State×Year	County + State×Year
Observations	12,428	11,155 / 12,197

No coefficient is significant at the 10% level. SEs clustered by county.

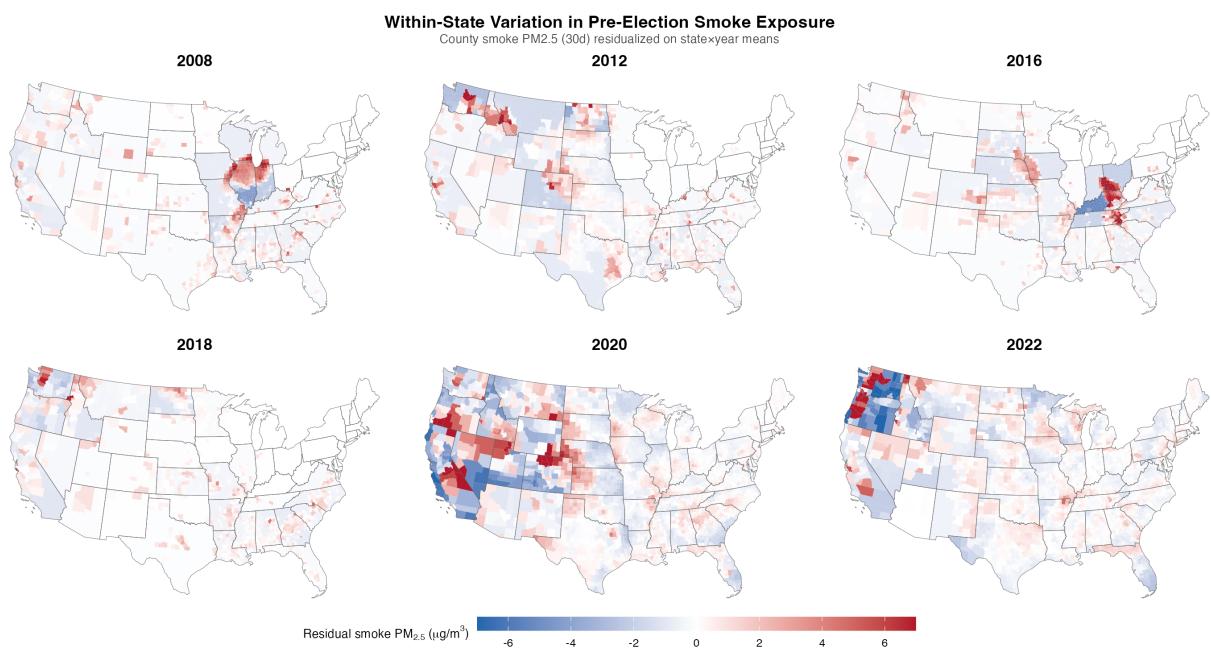


Figure A1: Within-state variation in pre-election smoke exposure. County-level smoke PM_{2.5} (30-day window) residualized on state-by-year means. The diverging scale shows counties with more (red) or less (blue) smoke than their state average. Meaningful within-state variation is concentrated in Western states during high-fire years.

C Literature Review

This appendix surveys the literatures most relevant to this paper, organized into five strands: environmental shocks and voting, the theory of anti-incumbent and retrospective voting, air pollution and political behavior, wildfire smoke and economic/health outcomes, and identification strategies using atmospheric variation.

C.1 Environmental Shocks, Natural Disasters, and Voting

A large literature examines whether environmental shocks alter political behavior. The findings cluster around two mechanisms: retrospective incumbent punishment and climate salience.

C.1.1 Incumbent Punishment

Healy and Malhotra (2010) provide the foundational analysis of myopic retrospective voting in response to disasters. Voters punish incumbents for disaster damage but reward them for disaster relief spending, and the asymmetry implies that governments underinvest in preparedness relative to relief. Achen and Bartels (2016) extend the argument, contending that voters engage in “blind retrospection,” punishing incumbents for events entirely beyond governmental control, including droughts, floods, and famously, shark attacks—though the shark attack finding is contested by Fowler and Hall (2018). Gasper and Reeves (2011) show that voters punish governors and presidents differently after natural disasters depending on whether a disaster declaration was issued, suggesting voters can partly distinguish between experiencing harm and receiving a government response. Bechtel and Hainmueller (2011) study the 2002 Elbe River flood in Germany and find that voters rewarded the incumbent government for effective flood management, demonstrating that the sign of the disaster–voting relationship depends on the government’s response. Cole et al. (2012) use rainfall variation in India to show that voters reward governments for disaster relief, providing a developing-country parallel.

C.1.2 Wildfires and Voting

Hazlett and Mildnerger (2020) is the most directly relevant paper. Using proximity to California wildfire perimeters, they find that fire proximity increases pro-environment ballot proposition voting, but only in already-Democratic areas. Republican areas show no response. My paper extends their work along three dimensions: smoke exposure rather than

fire proximity provides a plausibly exogenous treatment; national scope replaces a California-only sample; and I find effects across the partisan spectrum. Liao and Ruiz Junco (2022) study how extreme weather and natural disasters affect campaign contributions and House elections, finding that natural disasters hurt incumbents with anti-environment records, providing complementary evidence using campaign finance data.

C.1.3 Climate Salience

Herrnstadt and Muehlegger (2014) use Google search intensity to show that unusual weather increases searches for “climate change” and that members of Congress vote more pro-environment when their home states experience unusual weather—the clearest evidence for a salience channel linking environmental experience to political behavior. Hoffmann et al. (2022) use panel data across Europe to show that personal experience of climate-related events raises environmental concerns and increases Green party voting. Baccini and Leemann (2021) study Swiss referendum voting after floods and find that flood exposure increases support for pro-climate ballot measures by up to 20 percent. Elliott et al. (2023) examine whether U.S. senators vote more pro-environment after climate-related natural disasters, finding effects that are short-lived (about two years), consistent with the temporal dynamics in my results showing strongest effects at shorter pre-election windows.

C.1.4 Weather and Turnout

Gomez et al. (2007) show that rain on election day suppresses voter turnout with differential effects by party. This is relevant for the turnout channel: my positive turnout coefficient could be in tension with their finding if smoke operates like bad weather, but smoke exposure in my specification is measured over the pre-election window rather than on election day itself.

C.2 Anti-Incumbent Voting: Theoretical Foundations

The robust anti-incumbent finding places this paper within one of the oldest debates in democratic theory.

C.2.1 Sanctioning and Selection Models

The formal literature on electoral accountability begins with Ferejohn (1986), who models elections as a sanctioning mechanism: voters retain the incumbent if performance exceeds some threshold and replace her otherwise. The key insight is that retrospective voting

can sustain accountability even when voters are poorly informed about policy. Ashworth (2012) reviews how modern formal models blend sanctioning with selection (voters using past performance to infer incumbent quality), noting that the empirical literature has struggled to separate the two mechanisms.

C.2.2 Blind Retrospection and Its Critics

Achen and Bartels (2016) argue that retrospective voting is often “blind”—voters punish incumbents for events beyond governmental control. Gasper and Reeves (2011) push back, showing voters distinguish between harm and government response. Fowler and Hall (2018) challenge the shark attack result specifically. The emerging consensus, as articulated in Healy and Malhotra (2013), is a middle ground: voters are neither fully rational accountability agents nor purely blind avengers, but decision makers who sometimes apply coherent logic and sometimes fall prey to psychological biases.

C.2.3 Psychological Mechanisms

Three mechanisms are especially relevant. First, negative affect: Healy et al. (2010) show that irrelevant events (college football outcomes) shift incumbent evaluations, suggesting negative emotional states reduce support for the status quo regardless of source. Huber et al. (2012) provide experimental confirmation that voters overweight recent performance and are influenced by irrelevant lotteries. Second, attribution: partisanship powerfully shapes blame assignment after disasters, which may partly explain the partisan heterogeneity I find. Third, issue salience: negative experiences can shift the issues voters weight, activating voters who already care about climate while others simply experience negative affect. This dual-channel interpretation—salience for the pro-Democratic shift, affect for the anti-incumbent shift—maps onto my results where both channels appear to operate simultaneously through different pathways.

C.3 Air Pollution and Political Behavior

Bellani et al. (2024) is the closest analogue for the pollution–voting mechanism. Using 60 German federal and state elections, they exploit within-county variation in PM₁₀ on election day to show that a 10 $\mu\text{g}/\text{m}^3$ increase reduces the incumbent vote share by two percentage points. They argue this operates through a subconscious emotional channel: PM₁₀ fluctuations are imperceptible, yet they increase negative emotions. My setting differs critically in that wildfire smoke is visible and salient, potentially activating both their affect channel

and a deliberate salience channel. Kahn (2007) provides background for the heterogeneity pattern: areas with stronger baseline pro-environment preferences show larger smoke effects.

C.4 Wildfire Smoke: Economic, Health, and Behavioral Effects

A rapidly growing literature documents the broad consequences of wildfire smoke, establishing the health and mood channels that could mediate voting effects.

C.4.1 Labor Markets and Economic Costs

Borgschulte et al. (2022) is the leading paper using wildfire smoke as an exogenous source of PM_{2.5} variation to study economic outcomes, estimating that each additional smoke day reduces quarterly per capita earnings by about 0.1 percent. Their identification strategy—relying on the quasi-random spatial dispersion of smoke plumes using the same NOAA HMS data underlying Childs et al. (2022)—is closely related to mine. Importantly, they show that controlling flexibly for wind direction does not change their estimates, strengthening the case that smoke rather than correlated wind patterns drives the variation.

C.4.2 Mental Health and Mood

Burke et al. (2022) document behavioral and sentiment responses to wildfire smoke: during smoke events, people search more for air quality information, stay home more, and express more negative sentiment. This provides direct evidence for the mood mechanism I invoke for incumbent punishment. Jung et al. (2025) isolate short-term mental health effects of wildfire-specific PM_{2.5}, finding that a 10 µg/m³ increase significantly increases emergency department visits for depression and anxiety, with effects lasting up to seven days. Miller et al. (2024) estimate that wildfire smoke accounts for 18 percent of ambient PM_{2.5} concentrations and 0.42 percent of deaths among adults 65 and older, establishing the health harm of the exact pollution source I study.

C.5 Identification: Wind Direction and Atmospheric Dispersion

Deryugina et al. (2019) pioneered using wind direction as an instrument for PM_{2.5}, estimating effects of acute exposure on elderly mortality. The innovation is that the approach does not require knowing the location of pollution sources—it simply exploits the fact that wind direction shifts nonlocal pollution in and out of a county. My approach differs in that rather than using wind as an instrument for total PM_{2.5}, I use data from Childs et al. (2022) that has already separated wildfire smoke PM_{2.5} from other sources, and argue that

atmospheric dispersion is quasi-random conditional on county and year fixed effects. Rangel and Vogl (2019) use upwind/downwind variation from agricultural fires in Brazil to study birth outcomes, a conceptually very close design. Borgschulte et al. (2022) demonstrate as a robustness check that controlling for wind direction does not change their smoke estimates, a test that could be applied in my setting as well.